

CUSTOMER SEGMENTATION

Project Overview

Project Title: Customer Segmentation using Data Science Techniques.

Project Phase: Phase 3 – loading and preprocessing the customer data. Collect and preprocess the customer data for analysis.

Dataset Link:

<https://www.kaggle.com/datasets/akram24/mall-customers>

INTRODUCTION

In this Phase, We discussed about the Data Loading and the Data Preprocessing using the customer data and then also analysis the data. Loading and preprocessing customer data are fundamanetal steps in any data analysis.

Step 1: Data Loading

Data loading refers to the process of acquiring and importing data from external sources into a data storage system or a software application for further analysis, processing, or utilization. This data can come from various origins, including databases, files, web services, or APIs. The primary objective of data loading is to make the data available and accessible for use in tasks such as data analysis, reporting, machine learning, or any other data-driven processes.

Key characteristics and steps in data loading include:

1. Data Source Identification: Identifying the source of the data, which could be databases, flat files (e.g., CSV, Excel), web services, online platforms, or other data repositories.

2. Data Extraction: Extracting data from the source systems, which may involve querying databases, parsing files, or making web requests to retrieve data.

3. Data Transformation: Converting and restructuring the data as necessary to make it compatible with the target data storage or analysis platform. This may include data cleansing, data format conversion, and handling missing values.

4. Data Loading: Loading the transformed data into a data warehouse, data lake, database, or analysis tool for further processing. This often includes defining data schemas, mapping fields, and specifying data storage rules.

5. Data Validation: Ensuring the accuracy and integrity of the loaded data through various validation checks, such as data type verification, uniqueness checks, and referential integrity.

6. Data Indexing and Optimization: Creating indices and optimizing data storage structures to facilitate faster and more efficient data retrieval.

Data loading is a critical step in the data pipeline as it lays the foundation for data-driven decision-making and insights. Properly executed data loading ensures that data is readily available, accurate, and structured for analysis, reporting, and other applications, making it a cornerstone of data management and analytics processes.

Data loading and data preprocessing are essential steps in any data analysis or machine learning project. Here are some common methods and techniques used in these steps:

1. Using Libraries: Python provides several libraries for data loading, including:

- **`pandas`**: Used for reading data from various file formats like CSV, Excel, SQL, and more.
- **`numpy`**: Used for working with numerical data.
- **`openpyxl`** (for Excel files), **`sqlite3`** (for SQL databases), and others for specific data sources.

2. File Formats: Data can be stored in various formats, such as CSV, Excel, SQL databases, JSON, or even web scraping from HTML. You should choose the appropriate method for the data source.

3. Web APIs: For data from web services, you can use libraries like `requests` to make HTTP requests and fetch data from APIs.

Step 2: Data Preprocessing

1. Handling Missing Data:

- `dropna()`: Remove rows or columns with missing values.
- `fillna()`: Fill missing values with a specific value, like the mean, median, or a constant.
- Interpolation methods: Use statistical or time-based methods to estimate missing values.

```
ss.py ×
1 import numpy as np
2 import pandas as pd
3 from minisom import MiniSom
4 import matplotlib.pyplot as plt
5 df = pd.read_csv(r"C:\Users\dhana\OneDrive\Documents\IBM NM\Project Phase 1\Mall_Customers.csv")
6 # Load the Mall Customer dataset (replace 'file_path' with your actual dataset path)
7 # Check for missing data
8 missing_data = df.isnull().sum()
9 # Display the missing data count for each column
10 print("Missing Data:")
11 print(missing_data)
12 # Option 1: Remove rows with missing data
13 df_cleaned = df.dropna()
14 # Option 2: Fill missing data with the mean (or other strategy)
15 # Replace 'ColumnName' with the specific column name with missing data
16 # df['ColumnName'].fillna(df['ColumnName'].mean(), inplace=True)
17 # Display the cleaned dataset
18 print("\nCleaned Dataset:")
19 print(df_cleaned.head())
20
```

OUTPUT:

```
Run ss x
C:\python3\phase2\ven\Scripts\python.exe C:\python3\phase2\ss.py
Missing Data:
CustomerID      0
Genre           0
Age             0
Annual Income (k$)  0
Spending Score (1-100)  0
Unnamed: 5      200
Unnamed: 6      200
Total Spending   200
dtype: int64

Cleaned Dataset:
Empty DataFrame
Columns: [CustomerID, Genre, Age, Annual Income (k$), Spending Score (1-100), Unnamed: 5, Unnamed: 6, Total Spending]
Index: []

Process finished with exit code 0
```

EXPLANATION:

1. We load the Mall Customer dataset using pandas.
2. We check for missing data using the `isnull()` method, which returns a DataFrame of Boolean values (True for missing data, False for non-missing data).
3. We calculate and display the count of missing data for each column.
4. We provide two options for handling missing data:
 - Option 1: Removing rows with missing data using the `dropna()` method. This option is useful when you can afford to remove incomplete records.
 - Option 2: Filling missing data with the mean (or another strategy) using the `fillna()` method. You can uncomment this section and replace `'ColumnName'` with the specific column name with missing data.

Choose the option that best fits your data and analysis requirements. Handling missing data is essential to ensure that your analysis or machine learning models are based on complete and accurate data.

2. Encoding Categorical Data:

- `'One-Hot Encoding'`: Convert categorical variables into binary (0/1) columns for machine learning algorithms.
- `'Label Encoding'`: Assign unique integers to categorical values.

```
ss.py x
1 import numpy as np
2 import pandas as pd
3 from minisom import MiniSom
4 import matplotlib.pyplot as plt
5 df = pd.read_csv(r"C:\Users\dhana\OneDrive\Documents\IBM NM\Project Phase 1\Mall_Customers.csv")
6
7 # Let's assume 'Gender' is the categorical variable we want to encode
8 # Using one-hot encoding
9 data_encoded = pd.get_dummies(df, columns=['Genre'])
10
11 # Display the resulting DataFrame with one-hot encoding
12 print(data_encoded.head())
```

OUTPUT:

```
Run ss x
C:\python3\phase2\ven\Scripts\python.exe C:\python3\phase2\ss.py
  CustomerID  Age  ...  Genre_Female  Genre_Male
0         1   19  ...           False          True
1         2   21  ...           False          True
2         3   20  ...            True          False
3         4   23  ...            True          False
4         5   31  ...            True          False

[5 rows x 9 columns]

Process finished with exit code 0
```

3. Scaling and Normalization:

- 'MinMax Scaling': Scale numerical features to a specific range (e.g., [0, 1]).
- 'Standardization (Z-score scaling)': Scale numerical features to have a mean of 0 and a standard deviation of 1.

```
ss.py x
3 from minisom import MiniSom
4 import matplotlib.pyplot as plt
5 df = pd.read_csv(r"C:\Users\dhana\OneDrive\Documents\IBM NM\Project Phase 1\Mall_Customers.csv")
6 import pandas as pd
7 from sklearn.preprocessing import StandardScaler, MinMaxScaler
8
9
10 # Extract the numerical features to be scaled and normalized
11 numerical_features = df[['Age', 'Annual Income (k$)', 'Spending Score (1-100)']]
12
13 # Standardization (Z-score scaling)
14 scaler_standard = StandardScaler()
15 scaled_data_standard = scaler_standard.fit_transform(numerical_features)
16 data_standard = pd.DataFrame(scaled_data_standard, columns=numerical_features.columns)
17
18 # Min-Max Scaling
19 scaler_minmax = MinMaxScaler()
20 scaled_data_minmax = scaler_minmax.fit_transform(numerical_features)
21 data_minmax = pd.DataFrame(scaled_data_minmax, columns=numerical_features.columns)
22
23 # Display the scaled and normalized data
24 print("Standardized Data:")
25 print(data_standard.head())
26
27 print("\nMin-Max Scaled Data:")
28 print(data_minmax.head())
29
30
```

OUTPUT:

```
Run ss x
C:\python3\phase2\ven\Scripts\python.exe C:\python3\phase2\ss.py
Standardized Data:
   Age  Annual Income (k$)  Spending Score (1-100)
0 -1.424569         -1.738999          -1.723412
1 -1.281035         -1.738999          -1.706091
2 -1.352802         -1.700830          -1.688771
3 -1.137502         -1.700830          -1.671450
4 -0.563369         -1.662660          -1.654129

Min-Max Scaled Data:
   Age  Annual Income (k$)  Spending Score (1-100)
0  0.019231         0.000000         0.000000
1  0.057692         0.000000         0.005025
2  0.038462         0.008197         0.010050
3  0.096154         0.008197         0.015075
4  0.250000         0.016393         0.020101

Process finished with exit code 0
```

4. Feature Selection and Engineering:

- Select relevant features and remove irrelevant ones.
- Create new features based on domain knowledge or data analysis.

```
ss.py x
1
2 import pandas as pd
3 from minisom import MiniSom
4 import matplotlib.pyplot as plt
5 df = pd.read_csv(r"C:\Users\dhana\OneDrive\Documents\IBM NM\Project Phase 1\Mall_Customers.csv")
6
7 # Feature Engineering: Create a new feature "Total Spending" by adding "Annual Income" and "Spending Score"
8 df['Total Spending'] = df['Annual Income (k$)'] + df['Spending Score (1-100)']
9
10 # Display the first few rows of the dataset to verify the new feature
11 print(df.head())
12
13
14
```

OUTPUT:

```
Run ss x
C:\python3\phase2\ven\Scripts\python.exe C:\python3\phase2\ss.py
CustomerID  Genre  Age  ...  Unnamed: 5  Unnamed: 6  Total Spending
0          1   Male   19  ...         NaN         NaN             54
1          2   Male   21  ...         NaN         NaN             55
2          3  Female   20  ...         NaN         NaN             57
3          4  Female   23  ...         NaN         NaN             58
4          5  Female   31  ...         NaN         NaN             60

[5 rows x 8 columns]
```

5. Data Splitting:

- Split the dataset into training and testing sets for model evaluation.

```

ss.py x
1 import numpy as np
2 import pandas as pd
3 from minisom import MiniSom
4 import matplotlib.pyplot as plt
5 df = pd.read_csv(r"C:\Users\dhana\OneDrive\Documents\IBM NM\Project Phase 1\Mall_Customers.csv")
6 import pandas as pd
7 from sklearn.model_selection import train_test_split
8
9
10 # Split the data into features (X) and target (y)
11 X = df.drop(labels='Spending Score (1-100)', axis=1) # Features (excluding the target variable)
12 y = df['Spending Score (1-100)'] # Target variable
13
14 # Split the data into training and testing sets (adjust the test_size and random_state as needed)
15 X_train, X_test, y_train, y_test = train_test_split(*arrays: X, y, test_size=0.2, random_state=42)
16
17 # Display the shapes of the resulting sets to verify the split
18 print("Training set - X:", X_train.shape, " y:", y_train.shape)
19 print("Testing set - X:", X_test.shape, " y:", y_test.shape)
20

```

OUTPUT:

```

Run ss x
C:\python3\phase2\ven\Scripts\python.exe C:\python3\phase2\ss.py
Training set - X: (160, 7) y: (160,)
Testing set - X: (40, 7) y: (40,)
Process finished with exit code 0

```

6. Date and Time Handling:

- Extract date and time features from datetime columns.
- Convert datetime to numerical values for modeling.
- In this Dataset there is no date and time

If your dataset doesn't contain any date and time information, you cannot directly perform date and time handling on it. Date and time handling functions are designed to work with columns that contain date and time data. If your dataset doesn't have such data, there's no meaningful date and time handling to be done.

If you have other specific goals or analysis you want to perform on your dataset, please provide more details about your dataset and what you are trying to achieve. I'd be happy to help with alternative data processing or analysis tasks based on the actual data in your dataset.

7. Text Data Processing:

- Tokenization: Split text into words or phrases.
- Text cleaning: Removing special characters, stop words, and stemming or lemmatization.

```
ss.py x
1 import numpy as np
2 import pandas as pd
3 from minisom import MiniSom
4 import matplotlib.pyplot as plt
5 data = pd.read_csv(r"C:\Users\dhana\OneDrive\Documents\IBM NM\Project Phase 1\Mall_Customers.csv")
6 import pandas as pd
7 from sklearn.feature_extraction.text import CountVectorizer
8
9 # Select the 'Genre' column for text data processing
10 genre_data = data['Genre']
11
12 # Initialize a CountVectorizer to tokenize the text data
13 count_vectorizer = CountVectorizer()
14
15 # Tokenize the 'Genre' data
16 genre_matrix = count_vectorizer.fit_transform(genre_data)
17
18 # Convert the tokenized data to a DataFrame for better visualization (optional)
19 genre_df = pd.DataFrame(genre_matrix.toarray(), columns=count_vectorizer.get_feature_names_out())
20
21 # Display the tokenized data
22 print(genre_df)
23
```

OUTPUT:

```
Run  ss  x
C:\python3\phase2\ven\Scripts\python.exe C:\python3\phase2\ss.py
female  male
0       0    1
1       0    1
2       1    0
3       1    0
4       1    0
..      ...  ...
195     1    0
196     1    0
197     0    1
198     0    1
199     0    1

[200 rows x 2 columns]

Process finished with exit code 0
```

8. Outlier Detection and Handling:

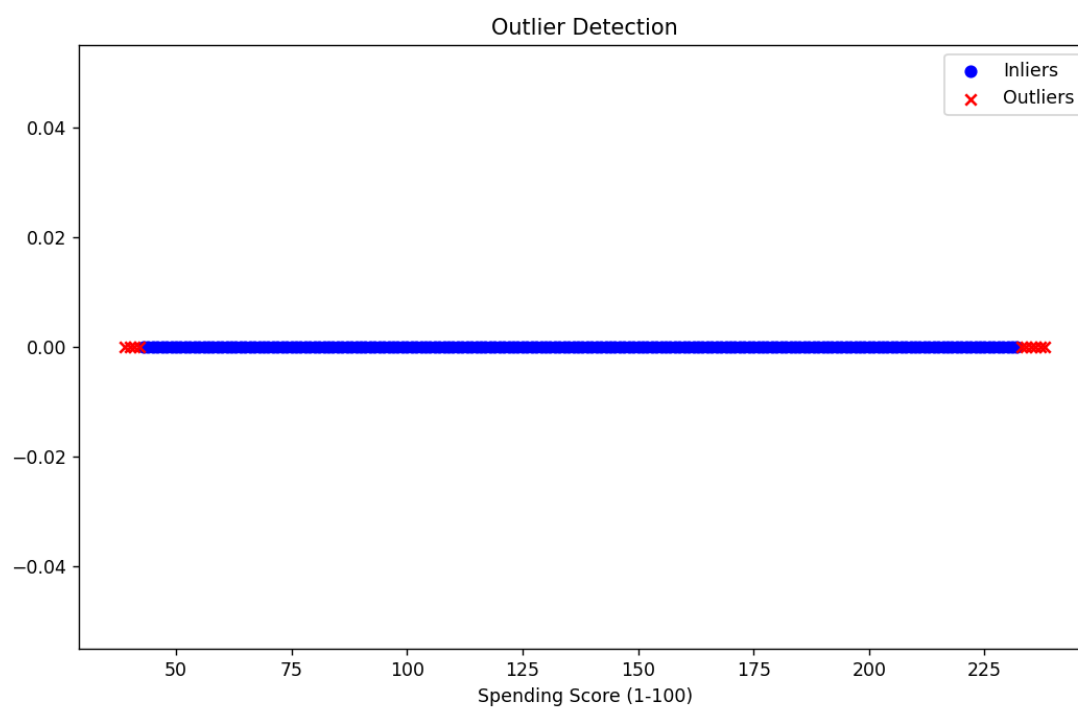
- Identify and deal with outliers that may affect the model's performance.

```

ss.py x
7 import numpy as np
8 from sklearn.ensemble import IsolationForest
9 import matplotlib.pyplot as plt
10
11
12 # Extract a numerical feature (e.g., 'Spending Score (1-100)' for this example)
13 feature_to_check = 'Spending Score (1-100)'
14 X = data[feature_to_check].values.reshape(-1, 1)
15
16 # Initialize and fit the Isolation Forest model for outlier detection
17 model = IsolationForest(contamination=0.05, random_state=42)
18 outlier_mask = model.fit_predict(X)
19
20 # Create a mask to identify outliers (1 for inliers, -1 for outliers)
21 is_inlier = outlier_mask == 1
22
23 # Separate outliers and inliers
24 outliers = X[~is_inlier]
25 inliers = X[is_inlier]
26
27 # Visualize the data with outliers highlighted
28 plt.figure(figsize=(10, 6))
29 plt.scatter(inliers, np.full_like(inliers, fill_value: 0), c='b', marker='o', label='Inliers')
30 plt.scatter(outliers, np.full_like(outliers, fill_value: 0), c='r', marker='x', label='Outliers')
31 plt.xlabel(feature_to_check)
32 plt.title('Outlier Detection')
33 plt.legend()
34 plt.show()
35

```

OUTPUT:



9. Normalization and Transformation:

- Apply mathematical transformations like log transformations to make data more suitable for modeling.

```
ss.py x
1 import numpy as np
2 import pandas as pd
3 from minisom import MiniSom
4 import matplotlib.pyplot as plt
5 data = pd.read_csv(r"C:\Users\dhana\OneDrive\Documents\IBM NM\Project Phase 1\Mall_Customers.csv")
6 import pandas as pd
7 from sklearn.preprocessing import StandardScaler
8 import numpy as np
9 # Assuming 'Annual Income (k$)' and 'Spending Score (1-100)' columns need normalization
10 columns_to_normalize = ['Annual Income (k$)', 'Spending Score (1-100)']
11 # Perform Standardization (Normalization) on the selected columns
12 scaler = StandardScaler()
13 data[columns_to_normalize] = scaler.fit_transform(data[columns_to_normalize])
14 # Perform Transformation (e.g., log transformation) on a column
15 data['Log_Transformed_Age'] = np.log(data['Age'])
16
17 # Display the first few rows of the transformed dataset
18 print(data.head())
19
```

OUTPUT:

```
Run ss x
C:\python3\phase2\ven\Scripts\python.exe C:\python3\phase2\ss.py
  CustomerID  Genre  Age  ...  Unnamed: 6  Total Spending  Log_Transformed_Age
0          1   Male   19  ...         NaN             NaN             2.944439
1          2   Male   21  ...         NaN             NaN             3.044522
2          3  Female   20  ...         NaN             NaN             2.995732
3          4  Female   23  ...         NaN             NaN             3.135494
4          5  Female   31  ...         NaN             NaN             3.433987

[5 rows x 9 columns]

Process finished with exit code 0
|
```

10. Data Visualization:


-Visualize the data to gain insights and detect patterns.

```

1 import numpy as np
2 import pandas as pd
3 from minisom import MiniSom
4 import matplotlib.pyplot as plt
5 data = pd.read_csv(r"C:\Users\dhana\OneDrive\Documents\IBM NM\Project Phase 1\Mall_Customers.csv")
6 import pandas as pd
7 import matplotlib.pyplot as plt
8 import seaborn as sns
9 # Display a pairplot to visualize pairwise relationships between numerical variables
10 sns.pairplot(data, diag_kind='kde', hue='Genre')
11 plt.title('Pairplot of Numerical Features')
12 plt.show()
13
14 # Create a histogram of Age distribution
15 plt.figure(figsize=(8, 6))
16 sns.histplot(data['Age'], bins=20, kde=True)
17 plt.title('Age Distribution')
18 plt.xlabel('Age')
19 plt.ylabel('Frequency')
20 plt.show()
21
22 # Create a bar plot to visualize the Gender distribution
23 plt.figure(figsize=(6, 4))
24 sns.countplot(data['Genre'])
25 plt.title('Genre Distribution')
26 plt.xlabel('Genre')
27 plt.ylabel('Count')
28 plt.show()

```

OUTPUT:



11. Scaling for Imbalanced Data:

- Techniques like oversampling or undersampling to address class imbalance in classification tasks.

```
ss.py x generic.py
1 import numpy as np
2 import pandas as pd
3 from minisom import MiniSom
4 import matplotlib.pyplot as plt
5 data = pd.read_csv(r"C:\Users\dhana\OneDrive\Documents\IBM NM\Project Phase 1\Mall_Customers.csv")
6 import pandas as pd
7 from sklearn.preprocessing import MinMaxScaler
8 # Select the 'Annual Income' column for scaling
9 income_data = data[['Annual Income (k$)']]
10 # Initialize the Min-Max scaler
11 scaler = MinMaxScaler()
12 # Fit and transform the data using Min-Max scaling
13 scaled_income = scaler.fit_transform(income_data)
14 # Create a new DataFrame with the scaled data
15 scaled_df = pd.DataFrame(scaled_income, columns=['Scaled Annual Income'])
16 # Concatenate the scaled data with the original DataFrame
17 data = pd.concat([data, scaled_df], axis=1)
18 # Display the first few rows of the updated DataFrame
19 print(data.head())
20 # Save the updated dataset if needed
21 data.to_csv('path_or_buf: 'scaled_mall_customer_dataset.csv', index=False) # Replace with your desired output filename
22
23
```

OUTPUT:

```
Run ss x
C:\python3\phase2\ven\Scripts\python.exe C:\python3\phase2\ss.py
CustomerID  Genre  Age  ...  Unnamed: 6  Total Spending  Scaled Annual Income
0           1   Male   19  ...         NaN           NaN           0.000000
1           2   Male   21  ...         NaN           NaN           0.000000
2           3  Female  20  ...         NaN           NaN           0.008197
3           4  Female  23  ...         NaN           NaN           0.008197
4           5  Female  31  ...         NaN           NaN           0.016393

[5 rows x 9 columns]

Process finished with exit code 0
```

12. Aggregation and Grouping:

-Summarize data by grouping it based on specific features.

```
ss.py x generic.py
2 import pandas as pd
3 from minisom import MiniSom
4 import matplotlib.pyplot as plt
5 df= pd.read_csv(r"C:\Users\dhana\OneDrive\Documents\IBM NM\Project Phase 1\Mall_Customers.csv")
6 import pandas as pd
7 # Grouping by Gender and finding the average spending for each group
8 gender_grouped = df.groupby('Genre')['Spending Score (1-100)'].mean()
9 # Grouping by Age and finding the total spending for each age group
10 age_grouped = df.groupby('Age')['Spending Score (1-100)'].sum()
11 # Aggregating data for each customer
12 agg_data = df.groupby('CustomerID').agg({
13     'Age': 'mean',
14     'Spending Score (1-100)': 'sum'
15 })
16 # Display the results
17 print("Average Spending by Gender:")
18 print(gender_grouped)
19 print("\nTotal Spending by Age:")
20 print(age_grouped)
21 print("\nAggregated Data for Each Customer:")
22 print(agg_data)
```

OUTPUT:

```
Run  ss  x
C:\python3\phase2\ven\Scripts\python.exe C:\python3\phase2\ss.py
Average Spending by Gender:
Genre
Female      135.562500
Male        142.238636
Name: Spending Score (1-100), dtype: float64

Total Spending by Age:
Age
18      459
19     1080
20      486
21      451
22      224
23      575
24      350
25      414
26      256
27      888
28      801
29      747
30     1194
31      904
32     2096
33      527
34      928
35      963
36     1139
37      467
```



```
Run ss x
⏮ ⏪ ⏩ ⏭ ⋮
↑
↓
↶
↷
📄
🗑
63      258
64       47
65      228
66      293
67      412
68      382
69       96
70      208
Name: Spending Score (1-100), dtype: int64

Aggregated Data for Each Customer:
      Age  Spending Score (1-100)
CustomerID
1        19.0             39
2        21.0             40
3        20.0             41
4        23.0             42
5        31.0             43
...        ...             ...
196       35.0            234
197       45.0            235
198       32.0            236
199       32.0            237
200       30.0            238

[200 rows x 2 columns]

Process finished with exit code 0
```